

Operating Systems Notes

0.1 Operating System Structure

1 Memory Management

Programs must be brought from disk into memory and then placed into a process.

The CPU can only access data from memory, not disk.

Register access takes at most 1 CPU clock, but Main Memory can take many cycles, causing a *Stall*.

1.1 Base and Limit Registers

A set of *base* and *limit registers* define the logical address space. the CPU must check every memory access is valid between the base and the limit for that user. Failure causes a trap to the OS monitor

1.2 Virtual Address Space

Logical/Virtual addresses are independent of physical memory.

Hardware translates virtual addresses into physical ones.

Logical/Virtual addresses a process can reference is called the address space.

1.3 Memory Management Unit (MMU)

Effectively is a hash function from logical address to physical address.

a MMU prevents the need for swapped out process to be swapped back into the same physical addresses.

Swapping is not typically supported on mobile devices, more likely to overwrite least used data.

1.4 Partitioning

Main memory is usually broken up into two partitions; The OS and user process.

Each process is contained within a single contiguous section on memory.

Reallocation registers are used to protect users processes from one another and from changing the OS code.

Some old techniques include:

- Fixed Partitions - simple but causes fragmentation often
- Variable Partitions - no internal fragmentation, but can leave holes in the physical memory

Dynamic Storage-Allocation is possible using First-fit, Best-fit and Worst-fit in terms of hole filling.

2 Paging

2.1 Address Translation Scheme

The Address generated by the CPU is divided into:

- Page Number - used as an index into the page table, containing the base address of each page in physical memory
- Page Offset - combined with the base address to define the physical memory address

2.2 Page Tables

Page Tables are kept in main memory.

The *Page-table base register (PTBR)* points to the page table

The *Page-table length register (PTLR)* indicates the size of the page table

In this scheme each instruction requires two memory accesses, one for the page table and one for the instruction, this can be solved by using *translation look-aside buffers (TLBs)*.

Some TLBs store address-space identifiers (*ASIDs*) in each TLB -

- uniquely identifies each process
- provides address-space protection

On a TLB miss, value is loaded into the TLB for faster access next time.

2.3 Memory Protection

A Valid/Invalid bit can be attached to each entry in a page table.

Valid indicates that the associated page is in the logical address space, and thus is a legal page.

We can use this or a *page-table length register (PTLR)* which can contain more info (Thanks for the detail here Michael).

Any violations result in a trap to the kernel.

2.4 Hierarchical Page Tables

Break up the logical address space into multiple page tables. A simple example is a two-leveled page table, known as *forward-mapped page table*.